

nn

December 6, 2019

```
In [3]: from keras.callbacks import ModelCheckpoint
        from keras.models import Sequential
        from keras.layers import Dense, Activation, Flatten
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_absolute_error
        from sklearn import metrics
        from matplotlib import pyplot as plt
        import pandas as pd
        import pickle

        data = None

        with open("../pickles/preprocessed_data.pkl","rb") as f:
            data = pickle.load(f)

        X = data.drop('log_price', 1)
        y = data['log_price']

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

Using TensorFlow backend.

In [4]: def build_NN(X_train, y_train, layer_width, layers, activation_fn, batch, loss_fn, epochs):
        NN_model = Sequential()

        # The Hidden Layers :
        NN_model.add(Dense(32, kernel_initializer='normal',input_dim = data.shape[1]-1, activation=activation_fn))
        for i in range(layers):
            NN_model.add(Dense(layer_width, kernel_initializer='normal',activation=activation_fn))

        # The Output Layer :
        NN_model.add(Dense(1, kernel_initializer='normal',activation='linear'))

        # Compile the network :
        NN_model.compile(loss = loss_fn, optimizer='adam', metrics=['mse','mae'])
        NN_model.summary()
```

```

prefix = str(layers + 1) + 'x' + str(layer_width) + '_' + activation_fn + '_' + lo
checkpoint_name = 'checkpoints/'+prefix+'_{val_loss:.5f}.hdf5' # Depth x width_l
checkpoint = ModelCheckpoint(checkpoint_name, monitor='val_loss', save_best_only =
callbacks_list = [checkpoint]

```

```

hist = NN_model.fit(X_train, y_train, epochs=epoch, batch_size=batch, verbose = 0,
return NN_model, hist, prefix

```

```

In [5]: def get_mse(model, X_test, y_test):
preds = NN_model.predict(X_test)
return metrics.mean_squared_error(preds,y_test)

```

```


In [6]: def show_plot(hist, name):
# plot loss during training
plt.title(name)
plt.plot(hist.history['loss'], label='train')
plt.plot(hist.history['val_loss'], label='val')
plt.legend()
plt.show()

```


```

In [5]: NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'relu', 32, 'mse') # width, d
show_plot(hist, name)
print(get_mse(NN_model, X_test, y_test))


```




0.15417947076354424




0.17042396764564222




0.15479445607862985




0.15829847452563972



0.15839866361927513



0.15302125004359599



0.15822405492642272


In [6]: # 32 width is optimal. 9 depth is optimal.

Checking for smaller widths


NN_model, hist, name = build_NN(X_train, y_train, 8, 4, 'relu', 32, 'mse') *# width, depth*

show_plot(hist, name)

print(get_mse(NN_model, X_test, y_test))



0.17307329965649632



0.1594820421067284


In [7]: # 32 width is optimal.

```
# checking for other parameters
```


```
NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'relu', 32, 'mae') # width, d
show_plot(hist, name)
print(get_mse(NN_model, X_test, y_test))
```

```
NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'tanh', 32, 'mse') # width, d
show_plot(hist, name)
print(get_mse(NN_model, X_test, y_test))
```

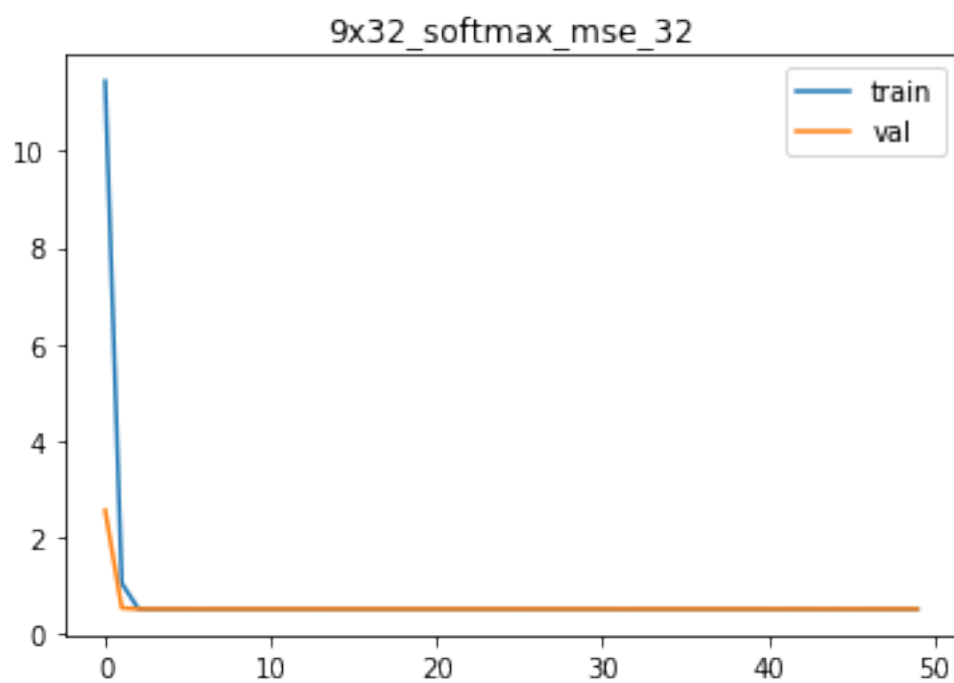
```
NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'softmax', 32, 'mse') # width
show_plot(hist, name)
print(get_mse(NN_model, X_test, y_test))
```



0.1581170385476567



0.5238097607726891




0.5158103607565511

```
In [8]: # Relu + linear is best. There's no overfitting.
```

```
    # Identifying best epochs
```

```
    NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'relu', 32, 'mse', 100) # width
    show_plot(hist, name)
```

```
    print(get_mse(NN_model, X_test, y_test))
```



0.17154473146846

```
In [9]: # 40 - 50 range seems to be the best epocs to avoid overfitting.
```

```
    # Final model.
```


```
    NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'relu', 32, 'mse', 50) # width
    show_plot(hist, name)
```

```
    print(get_mse(NN_model, X_test, y_test))
```

```
    preds = NN_model.predict(X_test)
```

```
    # Print r score
```

```
    print(metrics.r2_score(preds, y_test))
```



0.16519643511229784


0.5954635078887225

In [11]: *# Also trying with linea activation function*

```

NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'linear', 32, 'mse', 100) #
show_plot(hist, name)
print("mse: ", get_mse(NN_model, X_test, y_test))
preds = NN_model.predict(X_test)
# Print r score
print("rscore: ", metrics.r2_score(preds, y_test))

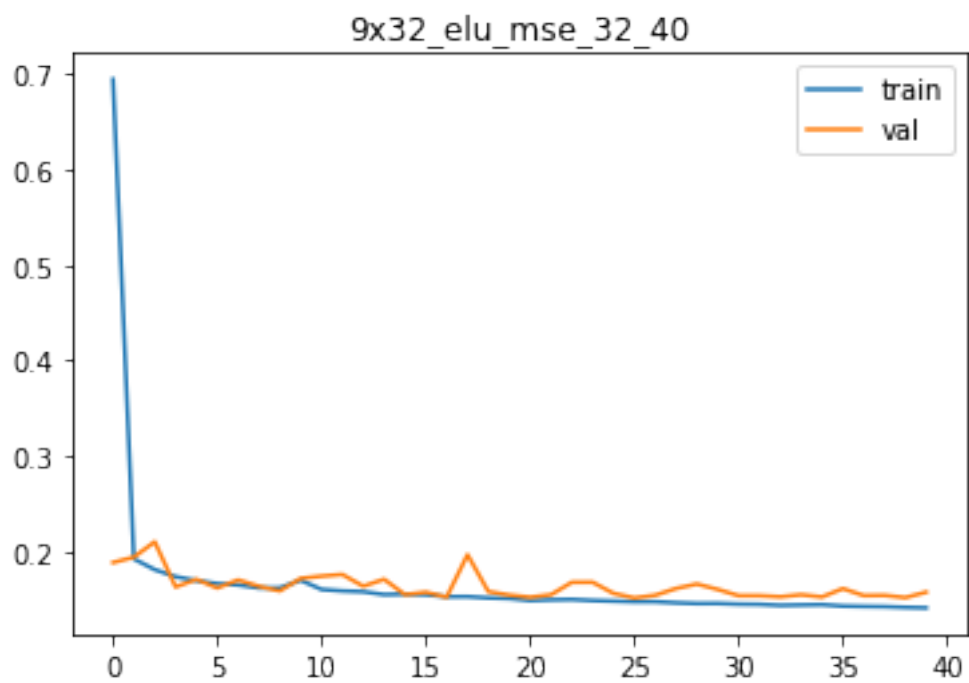
```



0.1651841126215015
0.583996829733165

In [14]: *# Also trying with exponential linear units (elu) activation function*

```
NN_model, hist, name = build_NN(X_train, y_train, 32, 8, 'elu', 32, 'mse', 40) # width
show_plot(hist, name)
print("mse: ", get_mse(NN_model, X_test, y_test))
preds = NN_model.predict(X_test)
# Print r score
print("rscore: ", metrics.r2_score(preds, y_test))
```



mse: 0.16766955781076806
rscore: 0.5644491352951815

In []: